

# Recommender Systems

This lecture is about recommender systems (or recommendation systems). In the meantime, we highlight the usefulness of Spark SQL, particularly when it relates to persistent tables.

## Spark SQL

As mentioned in the initial lectures, Spark SQL is a Spark module for structured data processing. It works alongside the APIs of DataFrame and Dataset and it is responsible for performing extra optimizations. We can also execute SQL queries and reading data from various files formats and Hive tables. (Apache Hive can manage large datasets residing in distributed storage using SQL)

Further details can be found in <https://spark.apache.org/docs/latest/sql-programming-guide.html> and <https://spark.apache.org/docs/latest/api/sql/index.html>

We can check the reference guide for Structured Query Language (SQL) which includes syntax, semantics, keywords, and examples for common SQL usage.

## Problem formulation

This exercise aims to build a recommender system of books, with focus on the recommendation model itself. The functional requirements for the Spark program we want to create are as follows:

1. To load the dataset and perform Exploratory Data Analysis (EDA), then store the information properly cleaned, including as SQL

tables.

2. To create a recommendation model supported by the ALS algorithm provided by Spark MLlib.
3. To pre-compute recommendations and store them in SQL tables.
4. To show recommendations.

## Dataset

The dataset describes 5-star rating and free-text tagging activity from MovieLens ( <https://movielens.org/> ), a movie recommendation service supported by the Social Computing Research Group at the University of Minnesota, USA ( <https://grouplens.org/> ). It contains 32000204 ratings and 2000072 tag applications across 87585 movies. These data were created by 200948 users between 9 January 1995 and 12 October 2023, and all selected users had rated at least 20 movies.

The data of concern are contained in the files `links.csv` , `movies.csv` , `ratings.csv` and `tags.csv` . The information included in the files are as follows:

- **links.csv**: Contains identifiers that link to sources of movie data e.g. to `movielens.org` or to `imdb.com` .
- **movies.csv**: Contains information about the movies, following the format `movieId, title, genres` .
- **ratings.csv**: Contains ratings of movies by users, following the format `userId, movieId, rating, timestamp` .
- **tags.csv**: Contains tags applied by users to movies, according to the format `userId, movieId, tag, timestamp` .

The given identifiers are common across the files.

The dataset (MovieLens 32M Dataset, file **ml-32m.zip**) can be downloaded from the location

`https://bigdata.iscte.me/abd/ml-32m.zip` .

## Initial settings

### Additional packages and imports

```
In [1]: import findspark, pyspark
```

```
from pyspark.sql import SparkSession
from pyspark.sql.types import *
import pyspark.sql.functions as F
```

In [2]: *# Some imports*

```
import os, sys

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

In [3]: *# Create the Spark session*

```
findspark.init()
findspark.find()

spark = SparkSession\
    .builder\
    .appName("MovieLens")\
    .config("spark.sql.shuffle.partitions",6)\
    .config("spark.sql.repl.eagereval.enabled",True)\
    .getOrCreate()
```

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

24/04/03 20:41:33 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j

In [4]: spark

Out [4]: **SparkSession - in-memory**

**SparkContext**

[Spark UI](#)

<b>Version</b>	v3.5.0
<b>Master</b>	local[*]
<b>AppName</b>	MovieLens

```
In [5]: # Some Spark related imports we will use hereafter
```

```
from pyspark.ml import Pipeline

from pyspark.ml.feature import StringIndexer
from pyspark.ml.recommendation import ALS
from pyspark.ml.evaluation import RegressionEvaluator
```

```
In [6]: from IPython.core.display import HTML
display(HTML("<style>pre { white-space: pre !important; }</style>"))
```

## Useful functions

```
In [7]: def plotBarColoured(df, xcol, ycol, colour):
        return sns.barplot(data=df, x=xcol, y=ycol, color=colour)
```

```
In [8]: def plotLine(df, xcol, ycol):
        return sns.lineplot(data=df, x=xcol, y=ycol)
```

```
In [9]: def plotHistogram(df, xcol, huecol=None):
        sns.histplot(data=df, x=xcol, hue=huecol, multiple="stack")
```

```
def plotScatter(df, xcol, ycol, huecol): sns.set_theme(style="white") sns.scatterplot(data=df, x=xcol, y=ycol, hue=huecol)def plotScatterMatrix(df, huecol): sns.pairplot(data=df, hue=huecol)
```

```
In [10]: def plotCorrelationMatrix(corr, annot=False):
# generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)
#cmap='coolwarm'

# draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0, annot=annot,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

```
In [11]: def plotBox(df, xcol, ycol, huecol=None, kind='box'):
return sns.catplot(data=df, x=xcol, y=ycol, hue=huecol, kind=kind)
```

```
In [12]: # Function to get columns of numeric type in a DataFrame
```

```
def numeric_columns(df):
    cls_numeric = []
    for x, t in df.dtypes:
        if t in ['int', 'double']:
            cls_numeric.append(x)
    return cls_numeric
```

```
In [13]: # Function to figure out the profile of nulls and uniques for columns in a DataFrame
```

```
def compute_nulls_and_uniques(df, cols):
    total = df.count()
    results = []
    for cl in cols:
        knulls = df.select(cl).filter(F.col(cl).isNull()).count()
        knullsporc = knulls / total
        knans = df.select(cl).filter(F.isnan(cl)).count()
        knansperc = knans / total
        kuniques = df.select(cl).distinct().count()
```

```
kuniquesperc = kuniques / total
results.append(Row(feature = cl, count_nulls = knulls, percentage_nulls = knullsporc,
                    count_nans = knans, percentage_nans = knansperc,
                    count_uniques = kuniques, percentage_uniques = kuniquesperc))

return spark.createDataFrame(results)
```

## Collect and label data

Checking working directory and data files

```
In [ ]: pwd
```

```
In [15]: data_dir =
```

```
In [ ]: ls -la
```

```
In [17]: ! head -n 3 ../Datasets/ml-32m/links.csv
```

```
movieId,imdbId,tmdbId
1,0114709,862
2,0113497,8844
```

```
In [18]: ! head -n 3 ../Datasets/ml-32m/movies.csv
```

```
movieId,title,genres
1,Toy Story (1995),Adventure|Animation|Children|Comedy|Fantasy
2,Jumanji (1995),Adventure|Children|Fantasy
```

```
In [19]: ! head -n 3 ../Datasets/ml-32m/ratings.csv
```

```
userId,movieId,rating,timestamp
1,17,4.0,944249077
1,25,1.0,944250228
```

```
In [20]: ! head -n 3 ../Datasets/ml-32m/tags.csv
```

```
userId,movieId,tag,timestamp  
22,26479,Kevin Kline,1583038886  
22,79592,misogyny,1581476297
```

## Reading the datasets

For the sake of this exercise, only data from the files `movies.csv` and `ratings.csv` are of interest.

```
In [21]: file_path = data_dir + "movies.csv"  
df_movies = spark.read.csv(file_path, header="true", inferSchema="true", sep=',')  
  
In [22]: file_path = data_dir + "ratings.csv"  
df_ratings = spark.read.csv(file_path, header="true", inferSchema="true", sep=',')
```

## Checking data

Schema, show and count.

### Movies

```
In [23]: df_movies.printSchema()  
df_movies.show(5, truncate=False)  
num_movies = df_movies.count()  
num_movies
```

```
root
|-- movieId: integer (nullable = true)
|-- title: string (nullable = true)
|-- genres: string (nullable = true)
```

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy

only showing top 5 rows

Out[23]: 87585

## Ratings

In [24]:



```
root
|-- userId: integer (nullable = true)
|-- movieId: integer (nullable = true)
|-- rating: double (nullable = true)
|-- timestamp: integer (nullable = true)
```

```
+-----+-----+-----+-----+
|userId|movieId|rating|timestamp|
+-----+-----+-----+-----+
|1      |17      |4.0   |944249077|
|1      |25      |1.0   |944250228|
|1      |29      |2.0   |943230976|
|1      |30      |5.0   |944249077|
|1      |32      |5.0   |943228858|
+-----+-----+-----+-----+
only showing top 5 rows
```

Out[24]: 32000204

## Explore and evaluate data

Let us get some data insight, with some **exploratory data analysis** based on descriptive statistics and visualizations.

## Datatypes

Is there a need to make adjustments/adding new fields to the data types specified in the dataframes?

Check the corresponding schema.

Issues:

- In Ratings, timestamp is set as integer but it is better using the Spark's timestamp datatype.

```
In [25]: # Let us create a new column but using just a simple cast. It may be enough!
```

```
df_ratings = ( df_ratings
                .withColumn(
                )
            )
```

```
In [26]: # Check the changes made
```

```
df_ratings.printSchema()
df_ratings.show(5, truncate=False)
```

root

```
|-- userId: integer (nullable = true)
|-- movieId: integer (nullable = true)
|-- rating: double (nullable = true)
|-- timestamp: integer (nullable = true)
|-- time: timestamp (nullable = true)
```

```
+-----+-----+-----+-----+-----+
|userId|movieId|rating|timestamp|time|
+-----+-----+-----+-----+-----+
|1      |17      |4.0   |944249077|1999-12-03 19:24:37|
|1      |25      |1.0   |944250228|1999-12-03 19:43:48|
|1      |29      |2.0   |943230976|1999-11-22 00:36:16|
|1      |30      |5.0   |944249077|1999-12-03 19:24:37|
|1      |32      |5.0   |943228858|1999-11-22 00:00:58|
+-----+-----+-----+-----+-----+
```

only showing top 5 rows

Check the README file provided, in particular in relation to movies and ratings data (timestamps, etc).

## Nulls, NaN and uniques

Identify number of nulls or NaN in columns as well uniques. This is helpful to further investigate the data.

```
In [27]: df_movies_nulls_uniques = compute_nulls_and_uniques(df_movies, df_movies.columns)
```

```
In [28]: df_movies_nulls_uniques.
```

feature	count_nulls	percentage_nulls	count_nans	percentage_nans	count_uniques	percentage_uniques
movieId	0	0.0	0	0.0	87585	1.0
title	0	0.0	0	0.0	87382	0.997682251527088
genres	0	0.0	0	0.0	1799	0.020540046811668666

```
In [29]: # cols_to_check = ['userID', 'movieId', 'rating', 'timestamp']
cols_to_check = ['userID', 'movieId', 'rating']
# but not including the new time column if we use the method below! Try.

df_ratings_nulls_uniques = compute_nulls_and_uniques(df_ratings, cols_to_check)
```

```
In [30]: df_ratings_nulls_uniques.
```

feature	count_nulls	percentage_nulls	count_nans	percentage_nans	count_uniques	percentage_uniques
userID	0	0.0	0	0.0	200948	0.006279584967645831
movieId	0	0.0	0	0.0	84432	0.0026384831796697297
rating	0	0.0	0	0.0	10	3.1249800782520014E-7

## Summary to figure out outliers

Summary of values for columns of interest. Use of describe() or summary()

```
In [31]: # df_movies.describe().show()
```

```
In [32]: # df_ratings.describe().show()
```

## Duplicates

```
In [33]: [num_movies, df_movies.dropDuplicates().count()]
```

```
Out[33]: [87585, 87585]
```


```
In [ ]: [num_ratings,
```

## Visualizations

Some visualizations to better understand the data. Feel free to adjust and/or add more.

```
In [35]: # Ratings distribution

df_plot = ( df_ratings.groupBy('rating')
             .count()
             .sort('count', ascending=False)
             .toPandas()
           )
plotBarColoured(df_plot, 'rating', 'count', 'orange')
plt.title('Counting of ratings')
plt.show()
```



```
In [36]: df_plot.head(40)
```

Out [36]:

	rating	count
0	4.0	8367654
1	3.0	6054990
2	5.0	4596577
3	3.5	4290105
4	4.5	2974000
5	2.0	2028622
6	2.5	1685386
7	1.0	946675
8	1.5	531063
9	0.5	525132

	rating	count
0	4.0	8367654
1	3.0	6054990
2	5.0	4596577
3	3.5	4290105
4	4.5	2974000
5	2.0	2028622
6	2.5	1685386
7	1.0	946675
8	1.5	531063
9	0.5	525132

In [37]: *# Top 10 popular movies based on number of ratings*

```
df_plot = ( df_ratings
            .groupBy('movieId')
            .agg({'userId': 'count', 'rating': 'average'})
            .sort('count(userId)', ascending=False)
            .limit(10)
            .toPandas()
            )
```

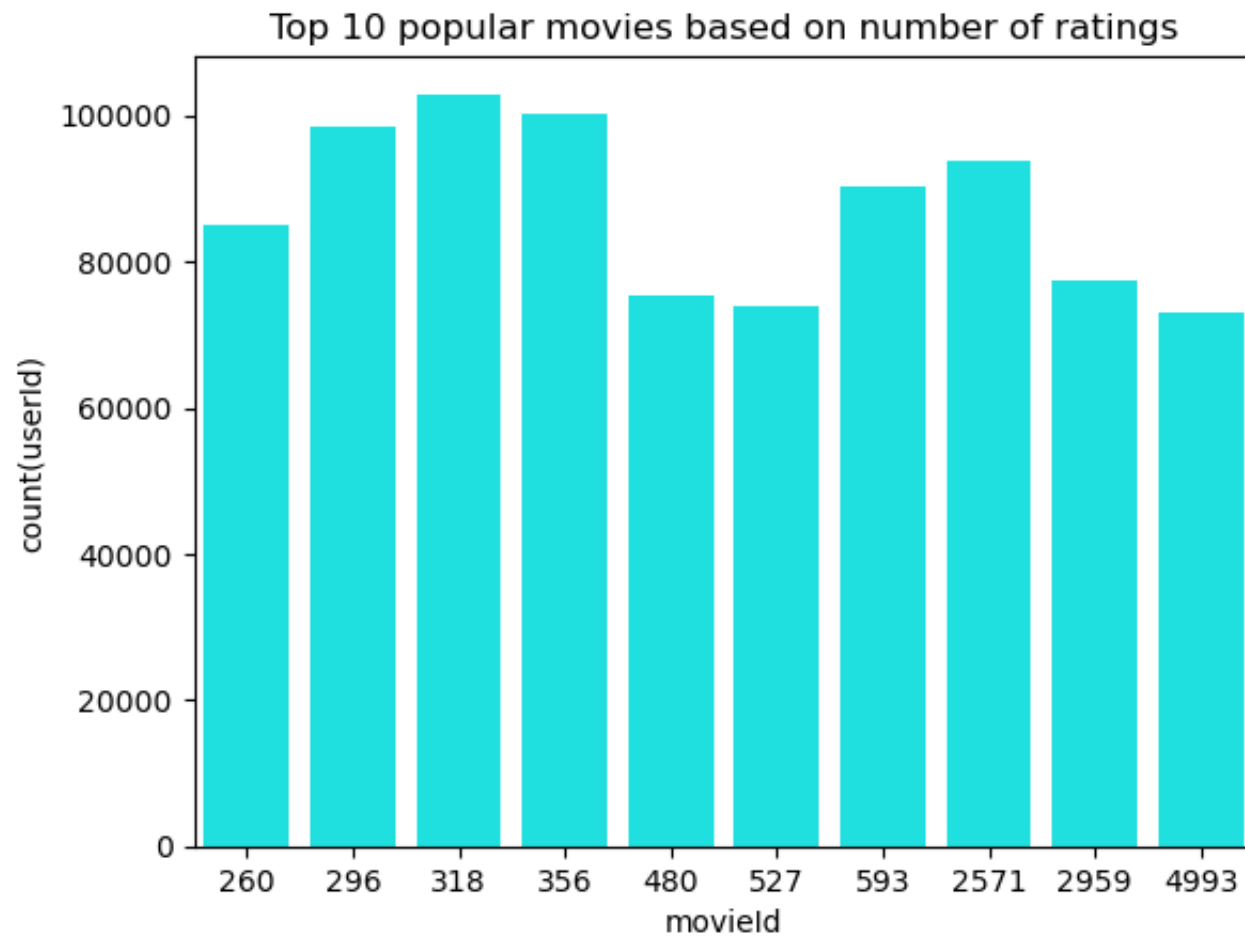
In [38]: df\_plot.

Out [38]:

	movieId	avg(rating)	count(userId)
--	---------	-------------	---------------

0	318	4.404614	102929
1	356	4.052744	100296
2	296	4.196969	98409
3	2571	4.156437	93808
4	593	4.148367	90330
5	260	4.099824	85010
6	2959	4.228780	77332
7	480	3.698623	75233
8	527	4.236990	73849
9	4993	4.092134	73122


```
In [39]: plt.title('Top 10 popular movies based on number of ratings')  
plt.show()
```



```
In [40]: plt.title('Average ratings of the top 10 popular movies based on number of ratings')  
plt.show()
```



Average ratings of the top 10 popular movies based on number of ratings



## Saving (clean) data

Saving data in proper format for further use, if needed.

**It seems data is OK.**

Context:

For a recommendation model, ratings data is critical.

As usual, we may want to have a smaller dataset just for the purpose of testing locally. That is, a smaller ratings dataset, but **keeping** the complete movies dataset, where movies are described.

```
In [41]: # from counting of ratings = 32000204

fraction = 0.3 # reduce to 30%

seed = 5
with_replacement = False
df_ratings_small = df_ratings.sample(withReplacement=with_replacement,
                                     fraction=fraction, seed=seed)
df_ratings_small.count()
```

Out[41]: 9597293

```
In [42]: # Delete memory consuming variables that are no longer needed, if any

# del ...
```

```
In [43]: # Saving movies in parquet format

output_movies = "movies.parquet"
df_movies.write.mode("overwrite").parquet(output_movies)
```

```
In [44]: # Saving ratings in parquet format

output_ratings = "ratings.parquet"
df_ratings.

output_ratings = "ratings_small.parquet"
df_ratings_small.
```

```
24/04/03 20:45:53 WARN MemoryManager: Total allocation exceeds 95.00% (1,020,054,720 bytes) of heap memory
Scaling row group sizes to 95.00% for 8 writers
24/04/03 20:46:05 WARN MemoryManager: Total allocation exceeds 95.00% (1,020,054,720 bytes) of heap memory
Scaling row group sizes to 95.00% for 8 writers
```

Check in the running directory if the operations above were properly accomplished.

Also, save them as persistent tables into Hive metastore.

Notes:

- An existing Hive deployment is not necessary to use this feature. Spark will take care of it.
- We can create a SQL table from a DataFrame with `createOrReplaceTempView` command, valid for the session. (there is also the option of global temporary views, to be shared among all sessions till the Spark application terminates)
- But with `saveAsTable`, there will be a pointer to the data in the Hive metastore. So persistent tables will exist even after the Spark program has restarted, as long as connection is maintained to the same metastore.

See details in <http://spark.apache.org/docs/latest/sql-data-sources.html>

```
In [45]: # Persistent tables into Hive metastore
```

```
df_movies.write.mode("overwrite").saveAsTable("MoviesTable")
df_ratings.write.mode("overwrite").saveAsTable("RatingsTable")
```

```
24/04/03 20:46:26 WARN MemoryManager: Total allocation exceeds 95.00% (1,020,054,720 bytes) of heap memory
Scaling row group sizes to 95.00% for 8 writers
```

## Feature engineering

### Data to be used

```
In [46]: # df_ratings_to_use = df_ratings
df_ratings_to_use = df_ratings_small
```

---

## Overview

After establishing the data to be used for the recommendation model, we should get an overview about what we have achieved, with some statistics and visualizations.

### But

we leave it as it is now, because there are no significant changes. Eventually, we could check the ratings and draw some plots, as it is the critical part of the system. You may try.

Nonetheless, let us recall the data types at stake, mainly because the ML algorithm requires numbers to process.

```
In [47]: df_ratings_to_use.printSchema()  
  
root  
|-- userId: integer (nullable = true)  
|-- movieId: integer (nullable = true)  
|-- rating: double (nullable = true)  
|-- timestamp: integer (nullable = true)  
|-- time: timestamp (nullable = true)
```

## Features transformation

```
In [48]: # Columns from ratings that are going to be considered in the model  
  
user_col = 'userId'  
item_col = 'movieId'  
rating_col = 'rating'
```

## Select and train model

In order to create the recommendation model, we will use the Alternating Least Squares (ALS) algorithm provided by Spark MLlib.

See details in <http://spark.apache.org/docs/latest/ml-collaborative-filtering.html> , as we advise to check the main assumptions the implemented algorithm relies upon. For example, notice that:

- it underlies a collaborative filtering strategy;
- it aims to fill in the missing entries of a user-item association matrix, in which users and items are described by a small set of latent factors that can be used to predict missing entries. The latent factors are learned by the ALS algorithm.

Again, as for data to train the model, the focus is on ratings.

## Partitioning of data

We will use the standard split 80/20, for the reasons explained in the lectures.

```
In [49]: # train/validation ratings split

df_train, df_validation = df_ratings_to_use.randomSplit([0.8, 0.2], 42)

# caching data ... but just the training part and if we want to (check the implications)
# df_train.cache()

print(f'There are {df_train.count()} rows in the training set and {df_validation.count()} in the validation set.
```

```
[Stage 131:=====> (2 + 6) / 8]
```

```
There are 7677972 rows in the training set and 1919321 in the validation set.
```

### Note:

As we did with clean data, we may consider storing the data split into files, should we want to use it elsewhere. This relates to the need of guaranteeing unicity in a different environment. We leave it as it is now.

# ALS model

Using the `ALS` estimator (the algorithm) to learn from the training data and consequently to build the model.

```
In [50]: # Build the recommendation model using ALS on the training data
# note that we set cold start strategy to 'drop' to ensure we don't get NaN evaluation metrics

als = ALS(maxIter=5,
          regParam=0.01,
          userCol=user_col,
          itemCol=item_col,
          ratingCol=rating_col,
          coldStartStrategy="drop",
          implicitPrefs=True
        )

# if the rating matrix is derived from another source of information
# (i.e. it is inferred from other signals), we may set implicitPrefs
# to True to get better results (see ALS reference)

# Also, other parameters that should be considered:
#     maxIter (maximum number of iterations),
#     rank (number of latent factors)
```

## ML pipeline configuration

```
In [51]: # The pipeline holds the only one stage set above

# As we will see below, we are going to use it just for evaluation purposes

pipeline =
```

## Model fitting

Get the model (as transformer) by fitting the pipeline to training data. It may take time!

```
In [52]: pipeline_model = pipeline.
```

```
24/04/03 20:47:33 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.blas.JNIBLAS
24/04/03 20:47:34 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.lapack.JNILAPACK
```

## Evaluate model

Let us evaluate the model that has been created.

## Testing the model

It is time to apply the model built to validation data. Again, we will use the pipeline set above.

Notice that, since the pipeline model is a transformer, we can easily apply it to validation data.

```
In [53]: # Make predictions on validation data and show values of columns of interest
```

```
df_prediction = pipeline_model.transform(df_validation)
```

```
In [54]: df_prediction.printSchema()
df_prediction.count()
```

root

```
|-- userId: integer (nullable = true)
|-- movieId: integer (nullable = true)
|-- rating: double (nullable = true)
|-- timestamp: integer (nullable = true)
|-- time: timestamp (nullable = true)
|-- prediction: float (nullable = false)
```

Out[54]: 1914660

```
In [55]: df_prediction.orderBy(user_col, item_col).show(truncate=False)
```

[Stage 360:>

(0 + 6) / 6]

userId	movieId	rating	timestamp	time	prediction
1	34	2.0	943228491	1999-11-21 23:54:51	0.20384896
1	166	5.0	943228442	1999-11-21 23:54:02	0.009173354
1	302	4.0	944253272	1999-12-03 20:34:32	0.036362004
1	835	3.0	944248888	1999-12-03 19:21:28	0.010275608
1	1120	1.0	944250121	1999-12-03 19:42:01	0.11650022
1	1203	5.0	944248888	1999-12-03 19:21:28	0.3472709
1	1276	3.0	944248943	1999-12-03 19:22:23	0.28005236
1	1810	3.0	944253272	1999-12-03 20:34:32	0.071456835
1	2918	4.0	943228400	1999-11-21 23:53:20	0.19806302
1	2966	1.0	943226846	1999-11-21 23:27:26	0.098449826
1	3088	3.0	944053949	1999-12-01 13:12:29	0.10106642
2	193	3.0	836423902	1996-07-03 20:58:22	0.102809034
2	236	4.0	836423512	1996-07-03 20:51:52	0.24210073
2	339	5.0	836423284	1996-07-03 20:48:04	0.4065569
2	508	5.0	836423970	1996-07-03 20:59:30	0.2735018
3	329	3.0	1084484653	2004-05-13 22:44:13	0.4210574
3	480	4.0	1084484525	2004-05-13 22:42:05	0.6463303
3	3033	3.0	1084485361	2004-05-13 22:56:01	0.19693887
3	5380	5.0	1084486046	2004-05-13 23:07:26	0.024898516
3	6889	5.0	1084484981	2004-05-13 22:49:41	0.0137242405

only showing top 20 rows

```
In [56]: df_prediction.orderBy(item_col, user_col).show(truncate=False)
```

[Stage 406:>

(0 + 6) / 6]



userId	movieId	rating	timestamp	time	prediction
176	1	2.5	1225231995	2008-10-28 22:13:15	0.15051374
237	1	2.5	1109494770	2005-02-27 08:59:30	0.07056041
249	1	5.0	852155678	1997-01-01 21:54:38	0.35842404
251	1	3.0	1497652791	2017-06-16 23:39:51	0.35108992
448	1	4.0	1084822895	2004-05-17 20:41:35	0.40689054
483	1	4.0	974933809	2000-11-22 22:56:49	0.48956427
686	1	3.0	1605147476	2020-11-12 02:17:56	-0.039479226
698	1	4.5	1289743558	2010-11-14 14:05:58	0.4781571
720	1	4.5	1440207472	2015-08-22 02:37:52	0.009225294
760	1	4.0	956636531	2000-04-25 05:22:11	0.0013743453
785	1	5.0	963726213	2000-07-16 06:43:33	0.4478437
881	1	5.0	1270793731	2010-04-09 07:15:31	0.8036039
978	1	5.0	1136487117	2006-01-05 18:51:57	0.57950664
997	1	3.0	1021669160	2002-05-17 21:59:20	-0.0027984977
1029	1	3.0	1485622552	2017-01-28 16:55:52	0.6691646
1043	1	4.0	1477711326	2016-10-29 04:22:06	0.55046403
1127	1	5.0	903199680	1998-08-15 17:48:00	0.63573647
1149	1	4.0	1672287580	2022-12-29 04:19:40	-0.0036828518
1173	1	4.0	877469440	1997-10-21 22:30:40	0.33966
1215	1	5.0	1652996991	2022-05-19 22:49:51	0.49569952

only showing top 20 rows

## Evaluation metrics

Let us use an evaluator.

```
In [57]: # Evaluate the model by computing the RMSE on the validation data
# We can use a different metric!
```

```
evaluator = RegressionEvaluator(metricName="rmse",
                                labelCol=rating_col,
```

```
predictionCol="prediction")

rmse = evaluator.evaluate(df_prediction)
print(f'The root mean square error is {rmse}.')
```

```
[Stage 452:> (0 + 6) / 6]
The root mean square error is 3.492908655744611.
```

Is it good?

## Saving the pipeline

```
In [58]: # We can save the pipeline for further use should we want

pipeline.

# later on, it can be loaded anywhere
```

Check in the working directory the pipeline that has been stored.

## Tune model

We can improve the model. For example, by changing parameters in the algorithm and also taking into consideration efficiency issues. We leave it now.

## Deploy model

### Pre-computing recommendations

The `ALS` algorithm provides some functions to get recommendations directly.

Although we can achieve results if working with predictions after the pipeline set (see below), we will take advantage of such

methods directly.

We should emphasize that, as it stands, we will not be using the pipeline for this task, but still, training the model as before.

```
In [59]: # Building the model with training data
```

```
model = als.fit(df_train)
```

```
In [60]: users = df_train.select(als.getUserCol()).distinct()  
movies = df_train.select(als.getItemCol()).distinct()
```

```
In [ ]: users.show()
```

```
In [ ]: movies.show()
```

```
In [ ]: [users.count(), movies.count()]
```

## Top movie recommendations for users (using PySpark's method)

```
In [65]: top_k_movies = 5  
recs_for_users = model.recommendForAllUsers(top_k_movies)
```

```
In [66]: recs_for_users.
```

```
[Stage 605:=====>(99 + 1) / 100]
```

	userId	recommendations
2		[[{356, 0.6291149}, {457, 0.5880901}, {364, 0.5095973}, {592, 0.4941127}, {380, 0.4807504}]]
4		[[{2571, 0.07897446}, {2762, 0.061630286}, {2959, 0.05767458}, {2706, 0.056905672}, {1527, 0.053370126}]]
5		[[{296, 0.3167363}, {318, 0.305773}, {480, 0.29812986}, {150, 0.29792345}, {110, 0.28963044}]]
10		[[{589, 0.755418}, {6539, 0.7338153}, {3793, 0.7308404}, {780, 0.7202}, {1210, 0.7128457}]]
14		[[{527, 0.03428884}, {3897, 0.02925688}, {1193, 0.02873656}, {4226, 0.028312184}, {2858, 0.028230254}]]
18		[[{527, 0.7010623}, {318, 0.696765}, {50, 0.6597826}, {593, 0.6515622}, {296, 0.6464614}]]
22		[[{260, 0.5152433}, {1196, 0.4729976}, {1210, 0.45681593}, {4993, 0.42755175}, {7153, 0.39240822}]]
25		[[{2571, 0.82313573}, {1198, 0.5778498}, {1196, 0.56281495}, {2959, 0.5405122}, {589, 0.5385877}]]
38		[[{260, 0.42504153}, {1210, 0.34767678}, {318, 0.320258}, {58559, 0.31045857}, {858, 0.31005093}]]
45		[[{2959, 0.26700395}, {593, 0.22511028}, {608, 0.21771216}, {50, 0.21588206}, {1, 0.21556452}]]
46		[[{79132, 0.56016713}, {356, 0.5497311}, {5952, 0.5050791}, {7153, 0.47048303}, {7361, 0.46800277}]]
50		[[{858, 0.59186554}, {260, 0.5414136}, {1221, 0.50605214}, {1193, 0.501761}, {318, 0.49956182}]]
59		[[{2858, 0.96588004}, {527, 0.9057325}, {1193, 0.7667044}, {912, 0.7149749}, {1198, 0.68696827}]]
60		[[{318, 0.54232097}, {858, 0.475461}, {296, 0.45072848}, {2028, 0.36494896}, {1221, 0.32939562}]]
73		[[{1, 0.4076958}, {593, 0.35796082}, {356, 0.3140781}, {608, 0.29410043}, {50, 0.28987113}]]
88		[[{260, 0.86008734}, {1196, 0.84301406}, {1198, 0.80369985}, {1210, 0.7872124}, {1270, 0.7434372}]]
90		[[{318, 0.40659332}, {527, 0.35766724}, {858, 0.31750238}, {260, 0.30590683}, {1210, 0.28576317}]]
97		[[{356, 0.51359725}, {5952, 0.5015791}, {593, 0.4903599}, {2571, 0.4634386}, {4993, 0.42179438}]]
102		[[{356, 0.26057163}, {593, 0.21973413}, {457, 0.21436344}, {589, 0.18578884}, {2858, 0.17432757}]]
108		[[{318, 0.4611783}, {858, 0.43999827}, {2028, 0.40568084}, {296, 0.32623303}, {2959, 0.29630783}]]

only showing top 20 rows

In [67]: `# Generate top movie recommendations for a specified set of users`

```
# subset_users = users.limit(5)
# recs_for_subset_users = model.recommendForUserSubset(subset_users, top_k_movies)
```

In [68]: `# recs_for_subset_users.show(truncate=False)`

Likewise, we can recommend users for each book. The model functions of concern are:

- `recommendForAllItems`

- *recommendForItemSubset*

Check the documentation. We may leave it as exercise.

## Movie recommendations for users (in general)

For the sake of running code in available time, we use only a subset of both users and movies.

```
In [ ]: m_users = 5
        n_movies = 10

        user_movie = users.limit(m_users).crossJoin(movies.limit(n_movies))
        user_movie.show()
        user_movie.count()
```

```
In [ ]: df_prediction = model.transform(user_movie).sort('userId', 'prediction', ascending=[True, False])
        df_prediction.show(truncate=False)
        df_prediction.count()
```

Note: The recommendations/predictions computed above (df\_prediction) may suggest movies that the users themselves have rated. We should make sure that those are filtered out prior to any movie recommendation. Try.

## Storing recommendations as persistent tables

Save the computed recommendations as persistent tables into the Hive metastore.

```
In [71]: recs_for_users.write.mode("overwrite").saveAsTable("RecommendationsForUsersTable")
```

```
In [ ]: df_prediction.
```

```
In [73]: # recs_for_subset_users.write.mode("overwrite").saveAsTable("RecommendationsForSubsetUsersTable")
```

Check the tables in the running directory.

## Exploring results

As a simple exercise, let us show the recommended books for a particular user, but using Spark SQL tables.

[illegible]

```
In [74]: # user to explore
```

```
user = 5
```

First, let us check the SQL tables.

```
In [75]: print(spark.catalog.listDatabases())
```

```
[Database(name='default', catalog='spark_catalog', description='default database', locationUri='file:/Users/adriar
```

```
In [76]: spark.catalog.listTables(dbName="default")
```

```
Out[76]: [Table(name='moviestable', catalog='spark_catalog', namespace=['default'], description=None, tableType='MANAGED',
Table(name='ratingstable', catalog='spark_catalog', namespace=['default'], description=None, tableType='MANAGED',
Table(name='rawrecommendationsforuserstable', catalog='spark_catalog', namespace=['default'], description=None,
Table(name='recommendationsforuserstable', catalog='spark_catalog', namespace=['default'], description=None, ta
```

```
In [77]: # Use managed tables
```

```
spark.sql("USE default")
```

```
Out[77]: DataFrame[]
```

```
In [78]: spark.catalog.listColumns('moviestable')
```

```
Out[78]: [Column(name='movieId', description=None, dataType='int', nullable=True, isPartition=False, isBucket=False),
Column(name='title', description=None, dataType='string', nullable=True, isPartition=False, isBucket=False),
Column(name='genres', description=None, dataType='string', nullable=True, isPartition=False, isBucket=False)]
```

```
In [79]: spark.catalog.listColumns('recommendationsforuserstable')
```

```
Out[79]: [Column(name='userId', description=None, dataType='int', nullable=True, isPartition=False, isBucket=False),
Column(name='recommendations', description=None, dataType='array<struct<movieId:int,rating:float>>', nullable=True, isPartition=False, isBucket=False)]
```

```
In [80]: spark.catalog.listColumns('rawrecommendationsforuserstable')
```

```
Out[80]: [Column(name='userId', description=None, dataType='int', nullable=True, isPartition=False, isBucket=False),
Column(name='movieId', description=None, dataType='int', nullable=True, isPartition=False, isBucket=False),
Column(name='prediction', description=None, dataType='float', nullable=True, isPartition=False, isBucket=False)]
```

```
In [81]: spark.sql("SELECT * FROM recommendationsforuserstable").show(10, truncate=False)
```

```
+-----+-----+
|userId|recommendations|
+-----+-----+
|13    |[{318, 0.7557082}, {110, 0.7321337}, {296, 0.68292063}, {527, 0.5774961}, {356, 0.5559001}]|
|37    |[{2571, 0.765461}, {2959, 0.6699872}, {593, 0.6166951}, {296, 0.58299637}, {1198, 0.57917607}]|
|49    |[{356, 0.5777449}, {150, 0.54368407}, {296, 0.5114955}, {318, 0.5094418}, {457, 0.50365126}]|
|52    |[{4993, 0.33177194}, {1136, 0.26097995}, {5618, 0.25628597}, {593, 0.24944158}, {4973, 0.24906562}]|
|55    |[{296, 0.09242468}, {318, 0.0850522}, {150, 0.08272076}, {590, 0.082091115}, {110, 0.08147114}]|
|57    |[{858, 0.41186503}, {1, 0.4062896}, {780, 0.37033325}, {733, 0.34334487}, {260, 0.33437893}]|
|83    |[{4993, 0.7257122}, {4306, 0.65895426}, {4886, 0.6567642}, {6377, 0.6546309}, {7153, 0.6428564}]|
|92    |[{1704, 0.32673088}, {356, 0.32437566}, {318, 0.22119159}, {1721, 0.22054507}, {3147, 0.21369512}]|
|107   |[{4993, 0.6715904}, {4886, 0.6625693}, {6377, 0.65656024}, {7153, 0.65485936}, {6539, 0.6478822}]|
|109   |[{364, 0.74043953}, {4306, 0.7198086}, {6377, 0.6720037}, {356, 0.6686075}, {4886, 0.63548297}]|
+-----+-----+
only showing top 10 rows
```

```
In [82]: spark.sql("SELECT * FROM moviestable").show(10, truncate=False)
```

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance
8	Tom and Huck (1995)	Adventure Children
9	Sudden Death (1995)	Action
10	GoldenEye (1995)	Action Adventure Thriller

only showing top 10 rows

```
In [83]: print(f'The recommended books for user {user} are: ')
```

The recommended books for user 5 are:

We leave it as exercise - the information needed is there!

## Additional exercises

1. Add a functionality to the current program so that, when the recommended books for a particular user are shown, it is also shown related information about the movies contained in the file `links.csv` mentioned earlier.

*Hint:* After reading the data from `links.csv`, and performing the subsequent checking, store the corresponding dataframe as a temporary view (see the method `createOrReplaceTempView`). Then, proceed similarly to the case of the table `moviestable`.

2. Given the current status of this notebook, redesign it so that the major tasks are split into various notebooks, or Python modules. The purpose is to modularize code having in mind the setup of a real recommender system. That is:
  - A downloader module, with focus on downloading data, cleansing it, and then storing it in a data store.



- A recommender building module, to create a recommendation model
- A recommender running module, to pre-compute recommendations and to save them in a data store.
- A recommender server, to retrieve recommendations upon queries made to the data store.

## References

- Learning Spark - Lightning-Fast Data Analytics, 2nd Ed. J. Damji, B. Wenig, T. Das, and D. Lee. O'Reilly, 2020
- <http://spark.apache.org/docs/latest/ml-guide.html>
- <https://docs.python.org/3/>
- F. Maxwell Harper and Joseph A. Konstan. 2015. The MovieLens Datasets: History and Context. ACM Transactions on Interactive Intelligent Systems (TiiS) 5, 4: 19:1–19:19.

In [ ]: